

# Data Exploration, Visualization, and Feature Engineering using R

*Yuhui Zhang, and Raja Iqbal*

## Basic plotting systems

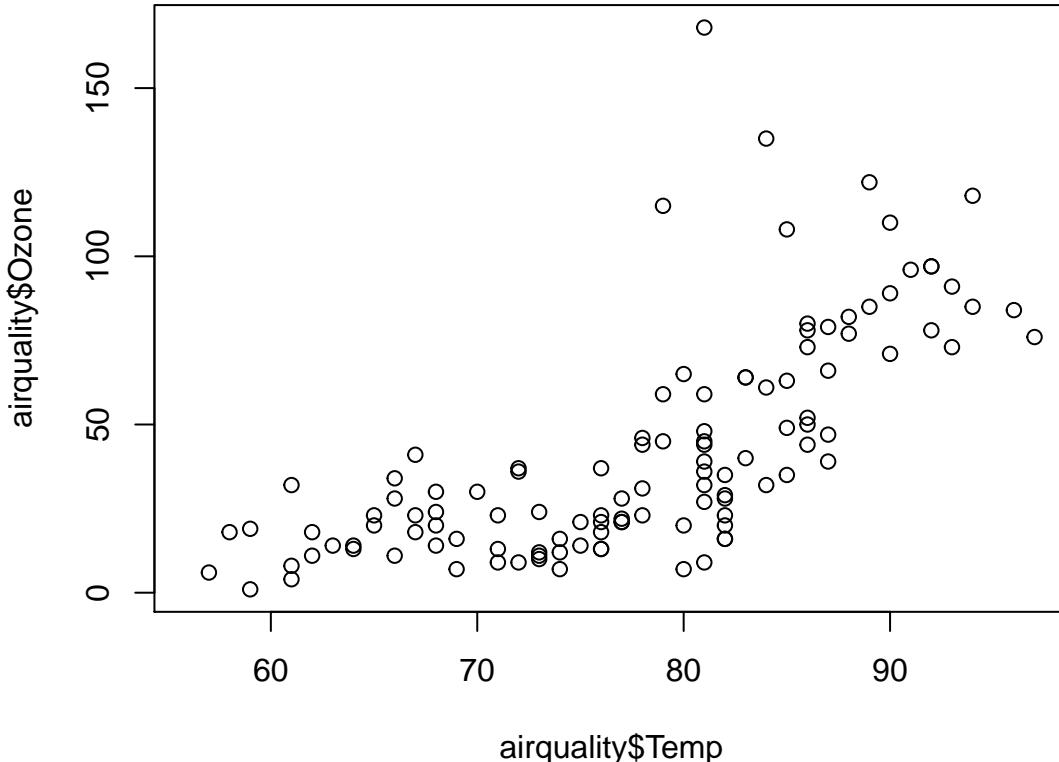
1. Base graphics: constructed piecemeal. Conceptually simpler and allows plotting to mirror the thought process.
2. Lattice graphics: entire plots created in a simple function call.
3. ggplot2 graphics: an implementation of the Grammar of Graphics by Leland Wilkinson. Combines concepts from both base and lattice graphics. (Need to install ggplot2 library)
4. Fancier and more telling ones.

A list of interactive visualization in R can be found at: <http://ouzor.github.io/blog/2014/11/21/interactive-visualizations.html>

---

### Base plotting system

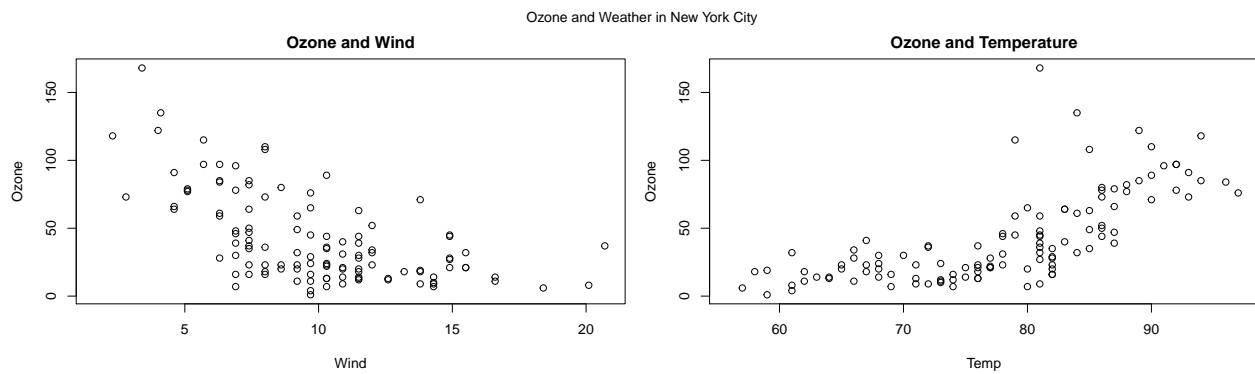
```
library(datasets)
## scatter plot
plot(x = airquality$Temp, y = airquality$Ozone)
```



---

## Base plotting system

```
## par() function is used to specify global graphics parameters that affect all plots in an R session.  
## Type ?par to see all parameters  
par(mfrow = c(1, 2), mar = c(4, 4, 2, 1), oma = c(0, 0, 2, 0))  
with(airquality, {  
  plot(Wind, Ozone, main="Ozone and Wind")  
  plot(Temp, Ozone, main="Ozone and Temperature")  
  mtext("Ozone and Weather in New York City", outer=TRUE)})
```



---

## Plotting functions (high level)

**PHASE ONE:** Mount a canvas panel on the easel, and draw the draft. (Initialize a plot.)

- `plot()`: one of the most frequently used plotting functions in R.
- `boxplot()`: a boxplot show the distribution of a vector. It is very useful to example the distribution of different variables.
- `barplot()`: create a bar plot with vertical or horizontal bars.
- `hist()`: compute a histogram of the given data values.
- `pie()`: draw a pie chart.

Remember to use `?plot` or `str(plot)`, etc. to check the arguments when you want to make more personalized plots. A tutorial of base plotting system with more details: <http://bcb.dfci.harvard.edu/~aedin/courses/BiocDec2011/2.Plotting.pdf>

---

## Plotting functions (low level)

**PHASE TWO:** Add more details on your canvas, and make an artwork. (Add more on an existing plot.)

- `lines`: adds lines to a plot, given a vector of x values and corresponding vector of y values

- points: adds a point to the plot
  - text: add text labels to a plot using specified x,y coordinates
  - title: add annotations to x,y axis labels, title, subtitles, outer margin
  - mtext: add arbitrary text to margins (inner or outer) of plot
  - axis: specify axis ticks
- 

## Save your artwork

R can generate graphics (of varying levels of quality) on almost any type of display or printing device. Like:

- postscript(): for printing on PostScript printers, or creating PostScript graphics files.
- pdf(): produces a PDF file, which can also be included into PDF files.
- jpeg(): produces a bitmap JPEG file, best used for image plots.

help(Devices) for a list of them all. Simple example:

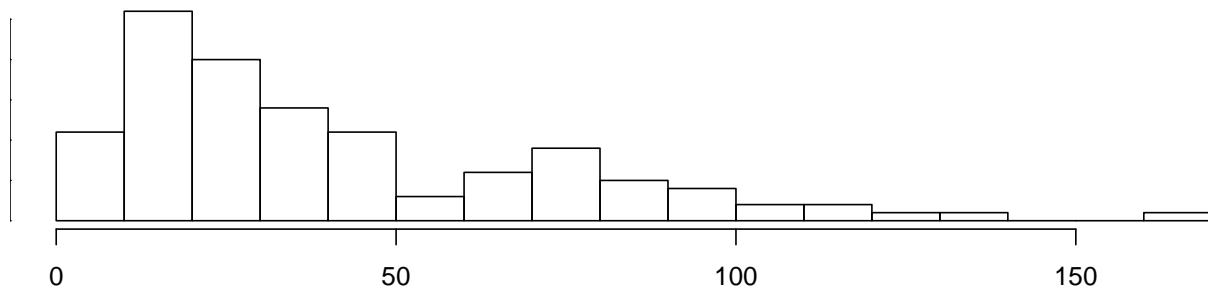
```
## png(filename = 'plot1.png', width = 480, height = 480, units = 'px')
## plot(x, y)
## dev.off()
```

---

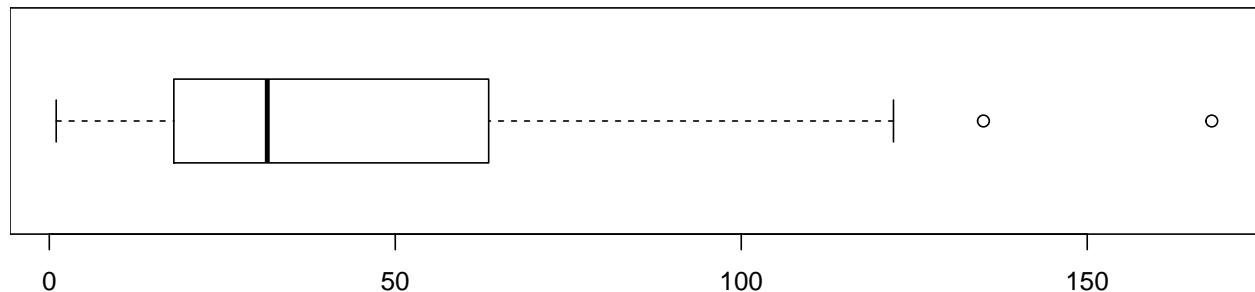
## Example: boxplot and histogram

```
## the layout
par(mfrow = c(2, 1), mar = c(2, 0, 2, 0), oma = c(0, 0, 0, 0))
## histogram at the top
hist(airquality$Ozone, breaks=12, main = "Histogram of Ozone")
## box plot below for comparison
boxplot(airquality$Ozone, horizontal=TRUE, main = "Box plot of Ozone")
```

### Histogram of Ozone

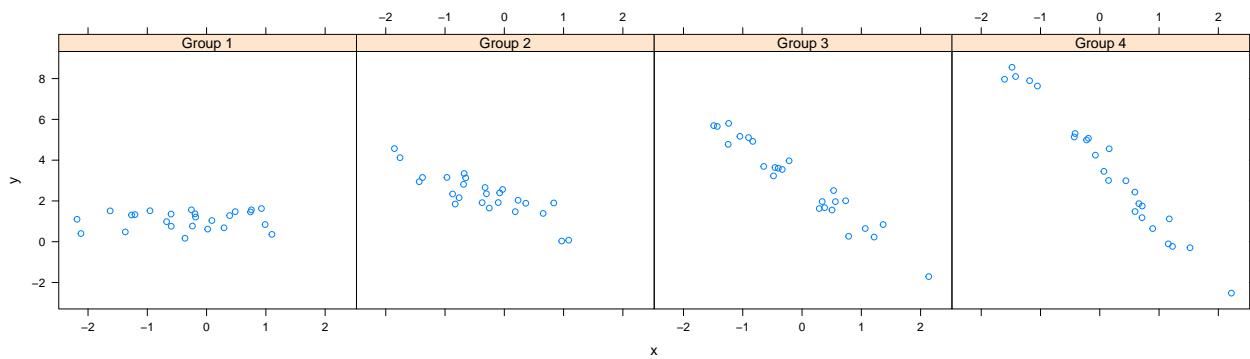


### Box plot of Ozone



### Lattice plotting system

```
library(lattice) # need to load the lattice library
set.seed(10) # set the seed so our plots are the same
x <- rnorm(100)
f <- rep(1:4, each = 25) # first 25 elements are 1, second 25 elements are 2, ...
y <- x + f - f * x + rnorm(100, sd = 0.5)
f <- factor(f, labels = c("Group 1", "Group 2", "Group 3", "Group 4"))
# first 25 elements are in Group 1, second 25 elements are in Group 2, ...
xyplot(y ~ x | f)
```



### Lattice plotting system

Want more on the plot? Customize the panel function:

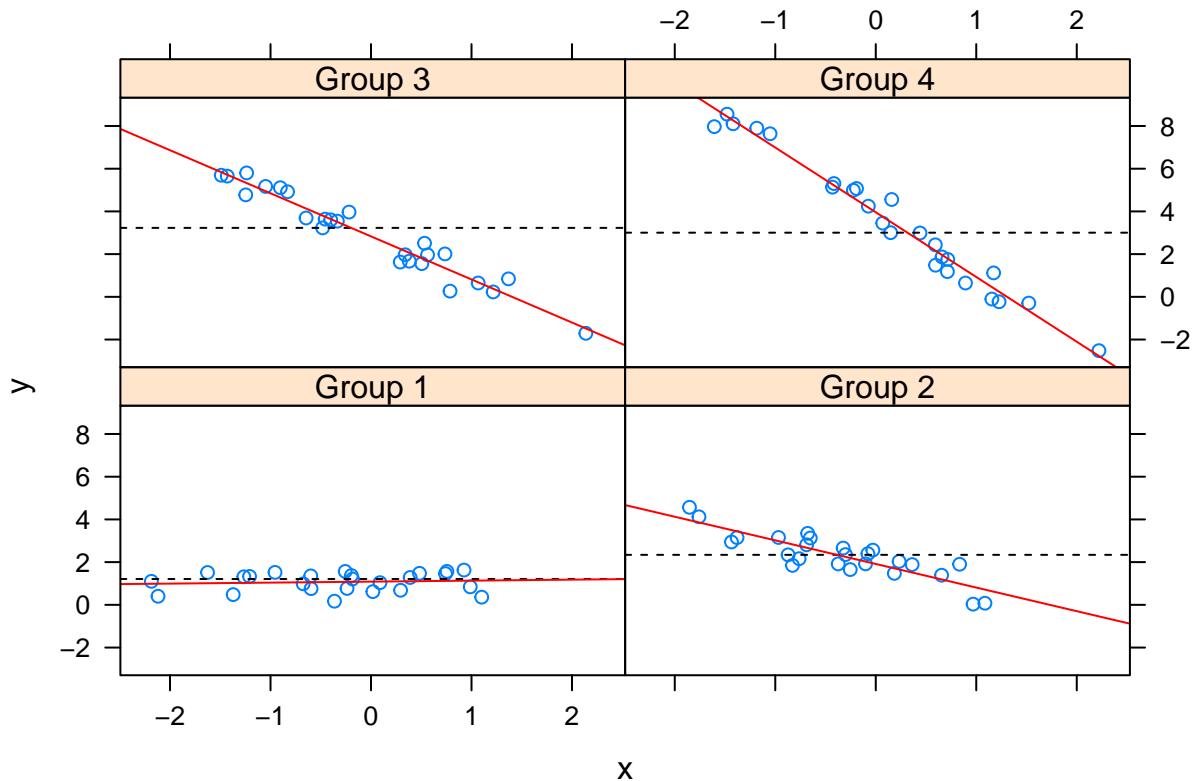
```

xyplot(y ~ x | f, panel = function(x, y, ...) {
  # call the default panel function for xyplot
  panel.xyplot(x, y, ...)
  # adds a horizontal line at the median
  panel.abline(h = median(y), lty = 2)
  # overlays a simple linear regression line
  panel.lmline(x, y, col = 2)
})

```

---

## Lattice plotting system



## Lattice plotting system

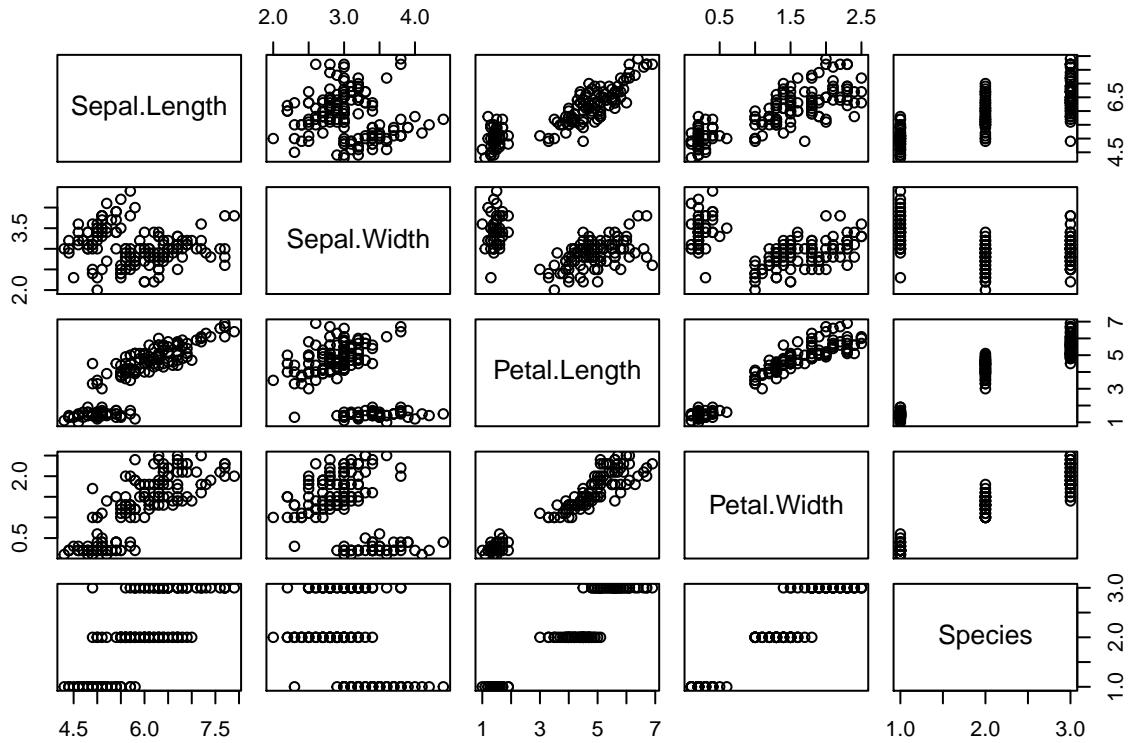
Plotting functions

- \* `xyplot()`: main function for creating scatterplots
- \* `bwplot()`: box and whiskers plots (box plots)
- \* `histogram()`: histograms
- \* `stripplot()`: box plot with actual points
- \* `dotplot()`: plot dots on “violin strings”
- \* `splom()`: scatterplot matrix (like `pairs()` in base plotting system)
- \* `levelplot() / contourplot()`: plotting image data

---

Very useful when we want a lot...

```
pairs(iris) ## iris is a data set in R
```



## ggplot2

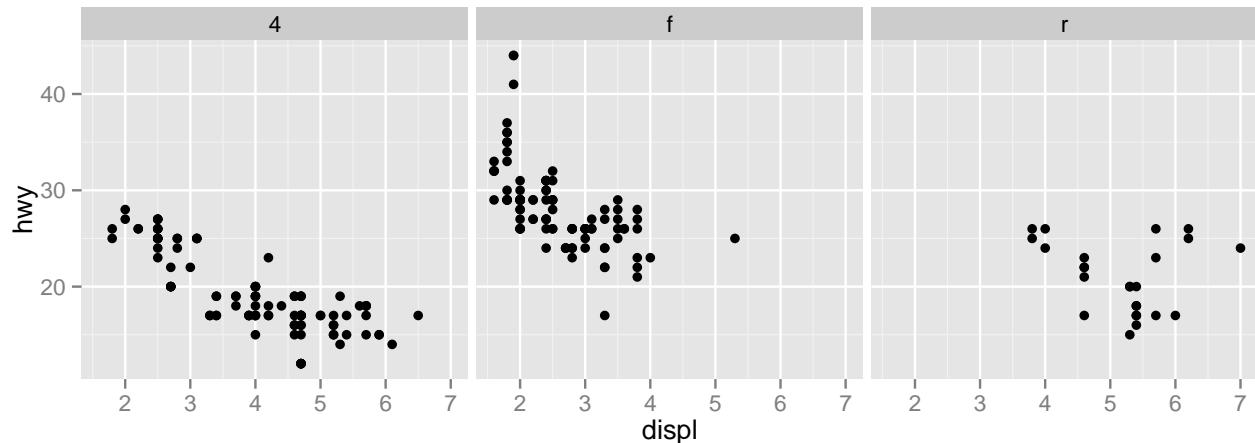
- An implementation of the Grammar of Graphics by Leland Wilkinson
- Written by Hadley Wickham (while he was a graduate student at Iowa State)
- A “third” graphics system for R (along with base and lattice)  
Available from CRAN via `install.packages()`  
web site: <http://ggplot2.org> (better documentation)
- Grammar of graphics represents the abstraction of graphics ideas/objects  
Think “verb”, “noun”, “adjective” for graphics  
“Shorten” the distance from mind to page
- Two main functions:  
`qplot()` hides what goes on underneath, which is okay for most operations `ggplot()` is the core function and very flexible for doing this `qplot()` cannot do

## qplot function

The `qplot()` function is the analog to `plot()` but with many build-in features  
Syntax somewhere in between base/lattice

Difficult to be customized (don’t bother, use full `ggplot2` power in that case)

```
library(ggplot2) ## need to install and load this library
qplot(displ, hwy, data = mpg, facets = .~drv)
```



---

## ggplot function

When building plots in ggplot2 (ggplot, rather than using qplot)

The “artist’s palette” model may be the closest analogy

Plots are built up in layers

Step I: Input the data

**noun:** the data

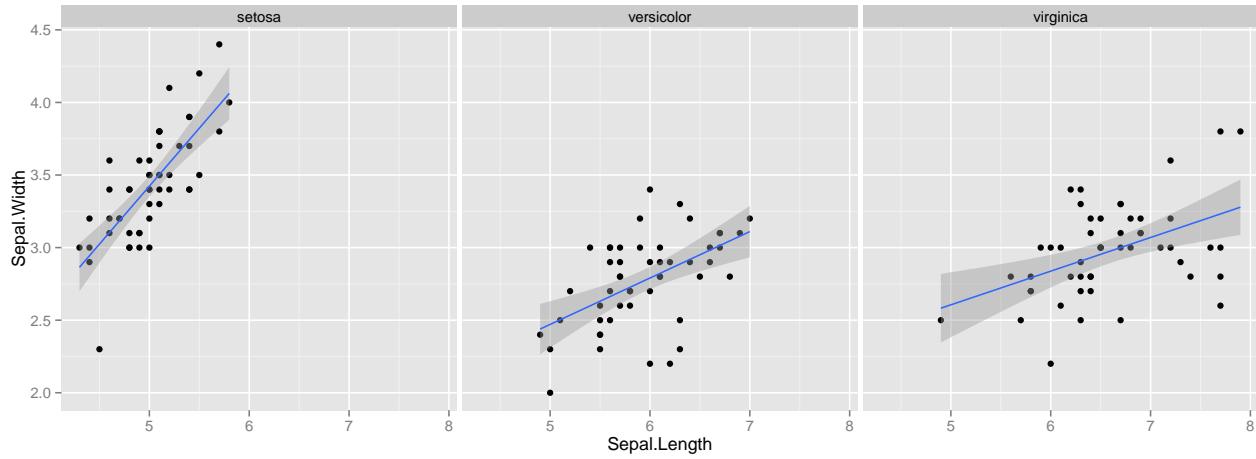
```
library(ggplot2) ## need to install and load this library
g <- ggplot(iris, aes(Sepal.Length, Sepal.Width)) ## this would not show you add plot
```

---

## ggplot function

- Step II: Add layers
- adjective:** describe the type of plot you will produce.

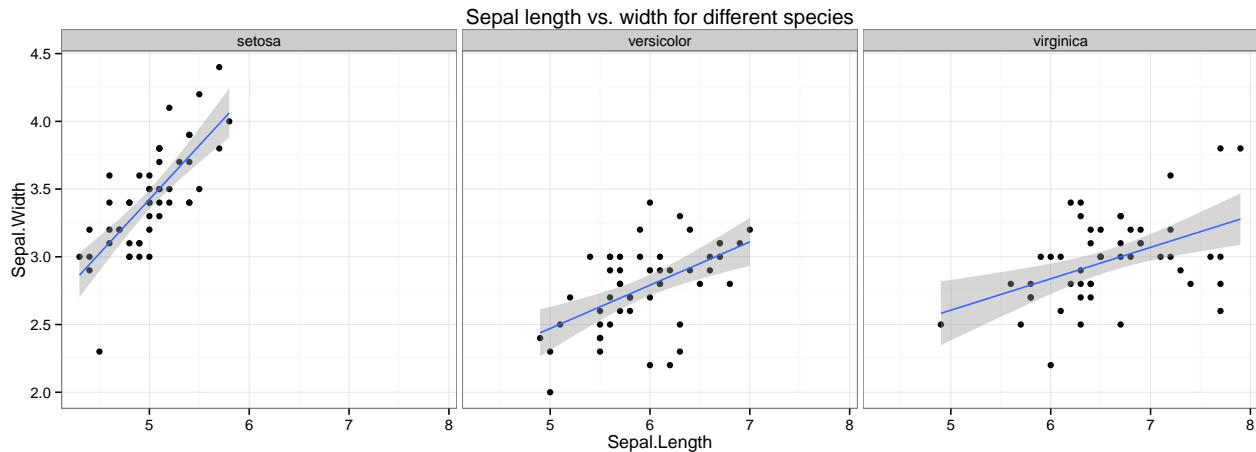
```
g + geom_point() + geom_smooth(method = "lm") + facet_grid(. ~ Species)
```



## ggplot function

- Step III: Add metadata and annotation  
**adjective**: control the mapping between data and aesthetics.

```
g <- g + geom_point() + geom_smooth(method = "lm") + facet_grid(. ~ Species)
g + ggtitle("Sepal length vs. width for different species") + theme_bw() ## verb
```



## Great documentation

Great **documentation** of ggplot with all functions in **step II** and **III** and demos:  
<http://docs.ggplot2.org/current/>

# Titanic tragedy data

---

## Reading RAW training data

- Download the data set “Titanic\_train.csv” from [https://raw.githubusercontent.com/datasciencedojo/datasets/master/Titanic\\_train.csv](https://raw.githubusercontent.com/datasciencedojo/datasets/master/Titanic_train.csv)
- Set working directory of R to the directory of the file using setwd()

```
titanic = read.csv('Titanic_train.csv')
```

---

## Look at the first few rows

What would be some good features to consider here?

```
options(width = 110)
head(titanic)
```

```
##   PassengerId Survived Pclass
## 1           1         0     3
## 2           2         1     1
## 3           3         1     3
## 4           4         1     1
## 5           5         0     3
## 6           6         0     3
## 
##             Ticket    Fare Cabin Embarked
## 1       A/5 21171 7.2500      S
## 2        PC 17599 71.2833     C85
## 3  STON/O2. 3101282 7.9250      S
## 4        113803 53.1000    C123
## 5        373450 8.0500      S
## 6        330877 8.4583      Q
## 
##           Name     Sex   Age SibSp Parch
## 1 Braund, Mr. Owen Harris   male  22      1
## 2 Cumings, Mrs. John Bradley (Florence Briggs Thayer) female 38      1
## 3 Heikkinen, Miss. Laina   female 26      0
## 4 Futrelle, Mrs. Jacques Heath (Lily May Peel) female 35      1
## 5 Allen, Mr. William Henry   male  35      0
## 6 Moran, Mr. James   male   NA      0
```

---

## What is the data type of each column?

```
sapply(titanic, class)
```

```
## PassengerId   Survived      Pclass        Name        Sex        Age      SibSp      Parch
##   "integer"   "integer"   "integer"   "factor"   "factor"   "numeric"   "integer"   "integer"
##   Fare        Cabin      Embarked
##   "numeric"   "factor"   "factor"
```

---

## Converting class label to a factor

```
titanic$Survived = factor(titanic$Survived, labels=c("died", "survived"))
titanic$Embarked = factor(titanic$Embarked, labels=c("unkown", "Cherbourg", "Queenstown", "Southampton"))
sapply(titanic, class)

## PassengerId      Survived       Pclass        Name         Sex          Age      SibSp      Parch
##    "integer"     "factor"     "integer"     "factor"     "factor"     "numeric"   "integer"   "integer"
##           Fare       Cabin Embarked
##    "numeric"     "factor"     "factor"

str(titanic$Survived)
##  Factor w/ 2 levels "died","survived": 1 2 2 2 1 1 1 1 2 2 ...
str(titanic$Sex)
##  Factor w/ 2 levels "female","male": 2 1 1 1 2 2 2 2 1 1 ...

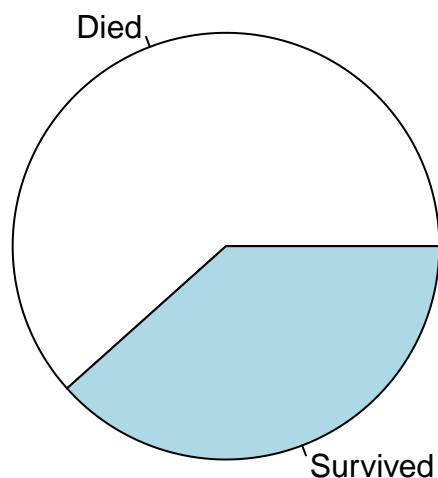
-----
```

## Class distribution - PIE Charts

```
survivedTable = table(titanic$Survived)
survivedTable

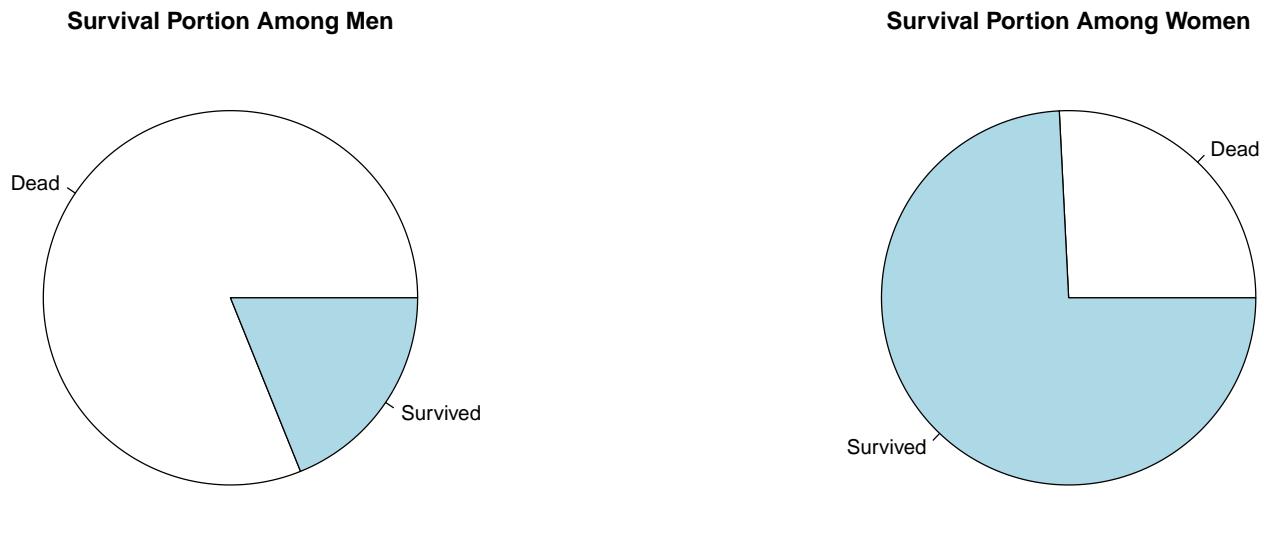
##
##      died survived
##      549      342

par(mar = c(0, 0, 0, 0), oma = c(0, 0, 0, 0))
pie(survivedTable, labels=c("Died", "Survived"))
```



## Is Sex a good predictor?

```
male = titanic[titanic$Sex=="male",]
female = titanic[titanic$Sex=="female",]
par(mfrow = c(1, 2), mar = c(0, 0, 2, 0), oma = c(0, 1, 0, 1))
pie(table(male$Survived), labels=c("Dead", "Survived"), main="Survival Portion Among Men")
pie(table(female$Survived), labels=c("Dead", "Survived"), main="Survival Portion Among Women")
```



## Is Age a good predictor?

```
Age <- titanic$Age; summary(Age)
```

```
##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.    NA's
##      0.42   20.12   28.00   29.70   38.00   80.00     177
```

How about summary segmented by **survival**

```
summary(titanic[titanic$Survived=="0",]$Age)
```

```
##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
##
```

```
summary(titanic[titanic$Survived=="1",]$Age)
```

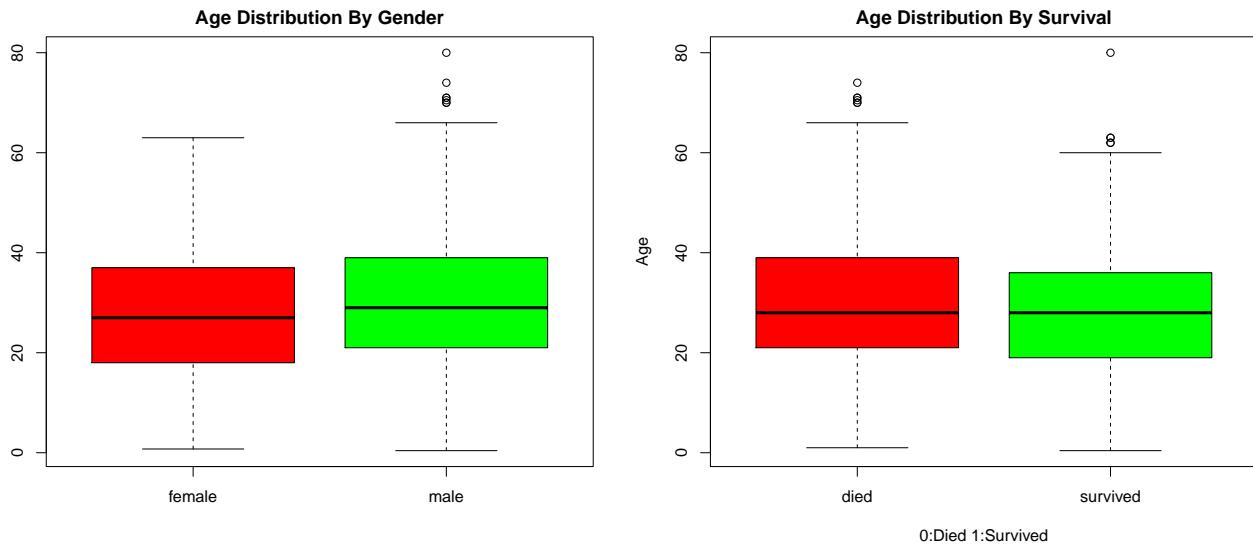
```
##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
##
```

## Age distribution by Survival and Sex

```

par(mfrow = c(1, 2), mar = c(4, 4, 2, 2), oma = c(1, 1, 1, 1))
boxplot(titanic$Age~titanic$Sex, main="Age Distribution By Gender", col=c("red","green"))
boxplot(titanic$Age~titanic$Survived, main="Age Distribution By Survival", col=c("red","green"),
        xlab="0:Died 1:Survived", ylab="Age")

```



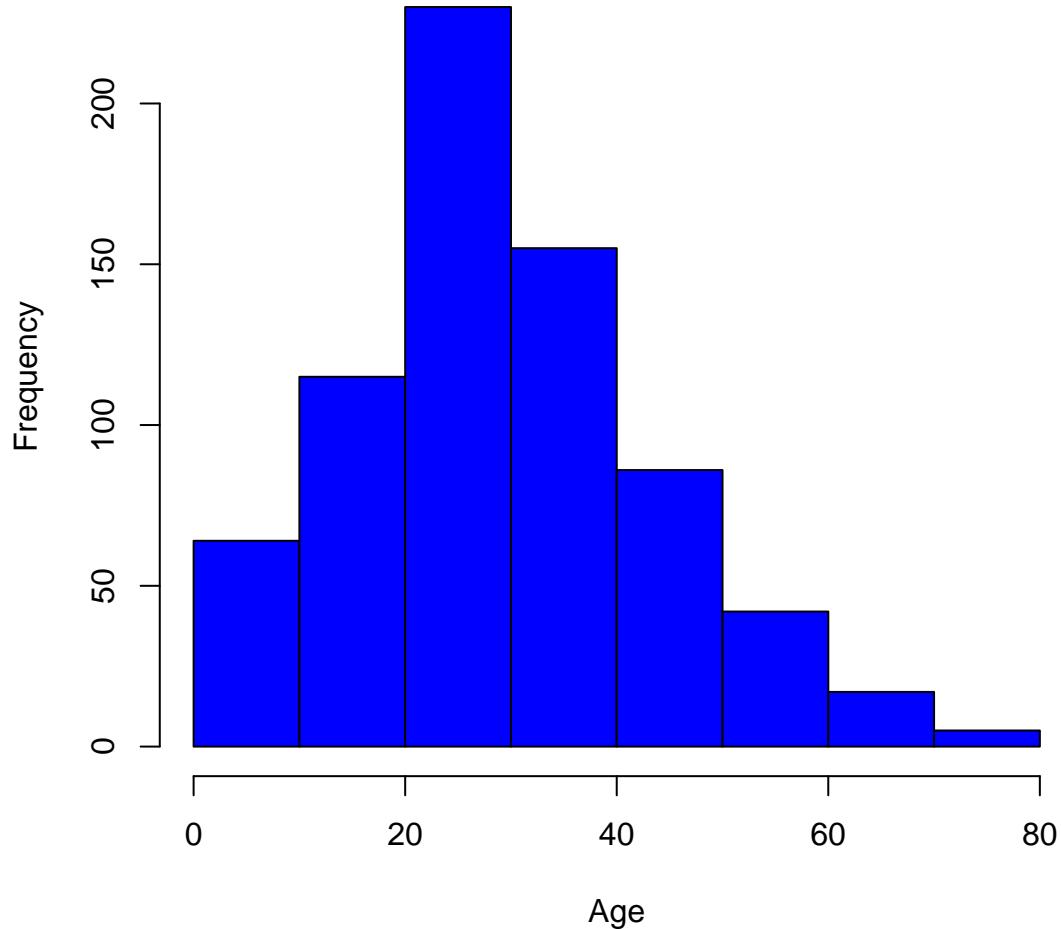
## Histogram of Age

```

hist(Age, col="blue", xlab="Age", ylab="Frequency",
     main = "Distribution of Passenger Ages on Titanic")

```

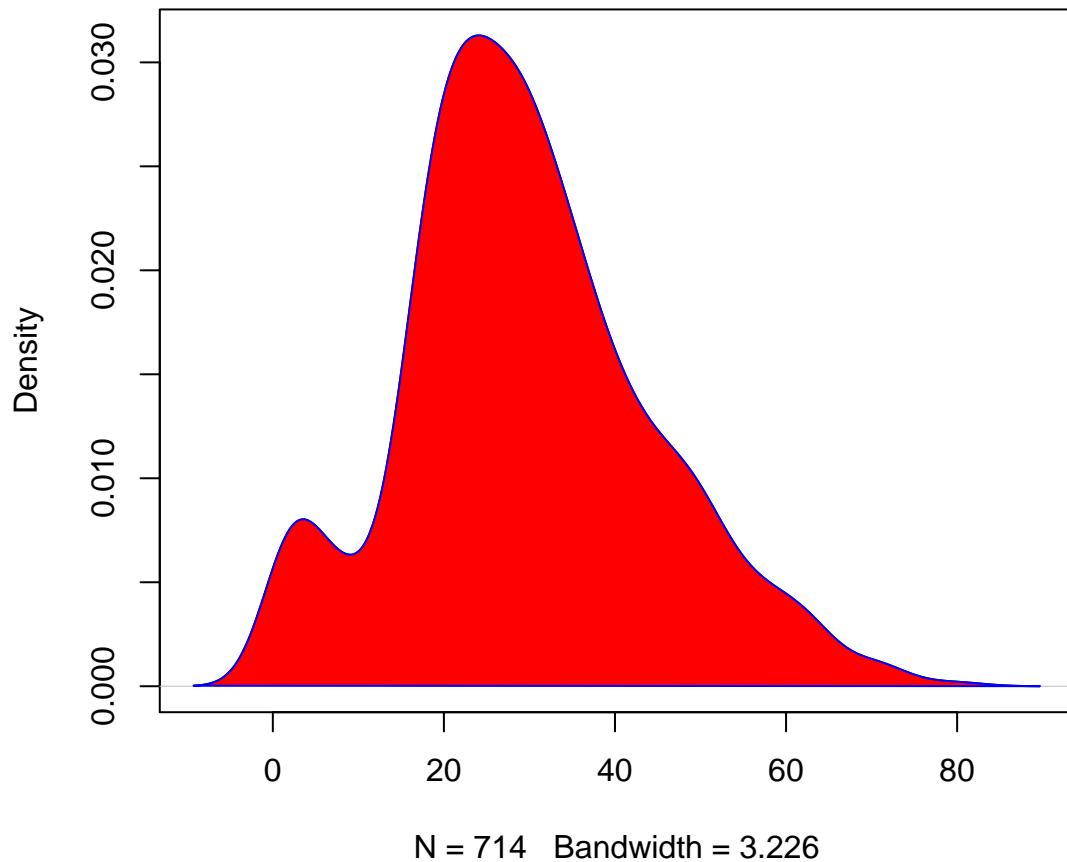
## Distribution of Passenger Ages on Titanic



Kernel density plot of age

```
d = density(na.omit(Age)) # density(Age) won't work, need to omit all NAs
plot(d, main = "kernel density of Ages of Titanic Passengers")
polygon(d, col="red", border="blue")
```

## kernel density of Ages of Titanic Passengers

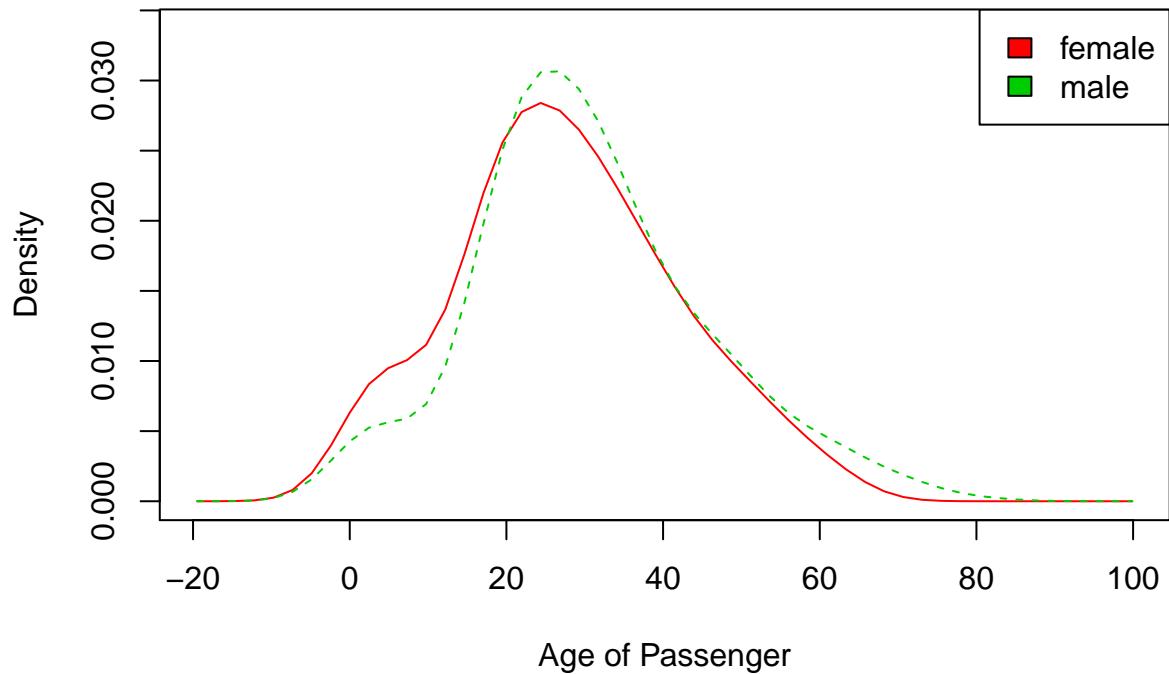


---

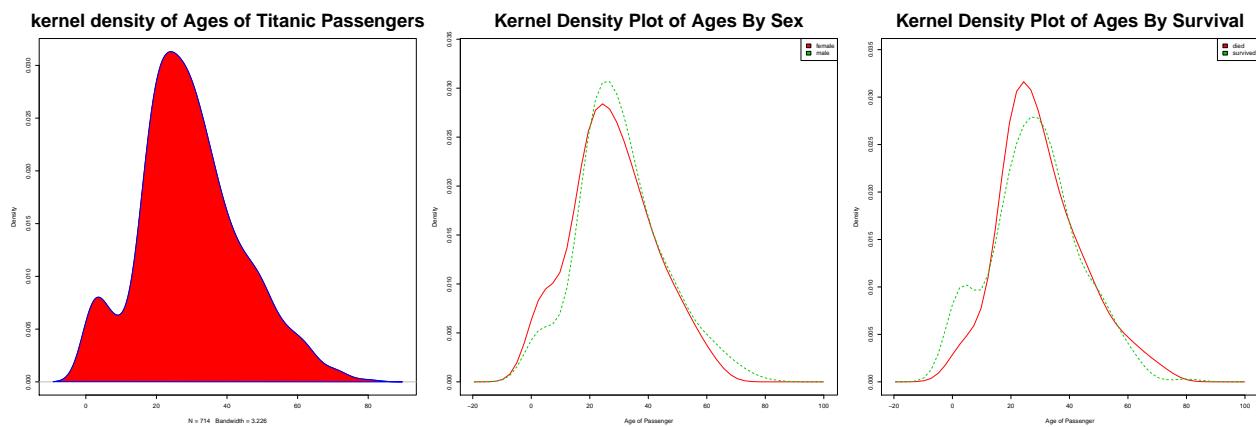
Comparison of density plots of Age with different Sex

```
## Package 'sm', version 2.2-5.4: type help(sm) for summary information
```

## Kernel Density Plot of Ages By Sex



Did Age have an impact on survival?



Create categorical groupings: Adult vs Child

An example of feature engineering!

```
## Multi dimensional comparison  
Child <- titanic$Age # Isolating age.
```

```

## Now we need to create categories: NA = Unknown, 1 = Child, 2 = Adult
## Every age below 13 (exclusive) is classified into age group 1
Child[Child<13] <- 1
## Every child 13 or above is classified into age group 2
Child[Child>=13] <- 2

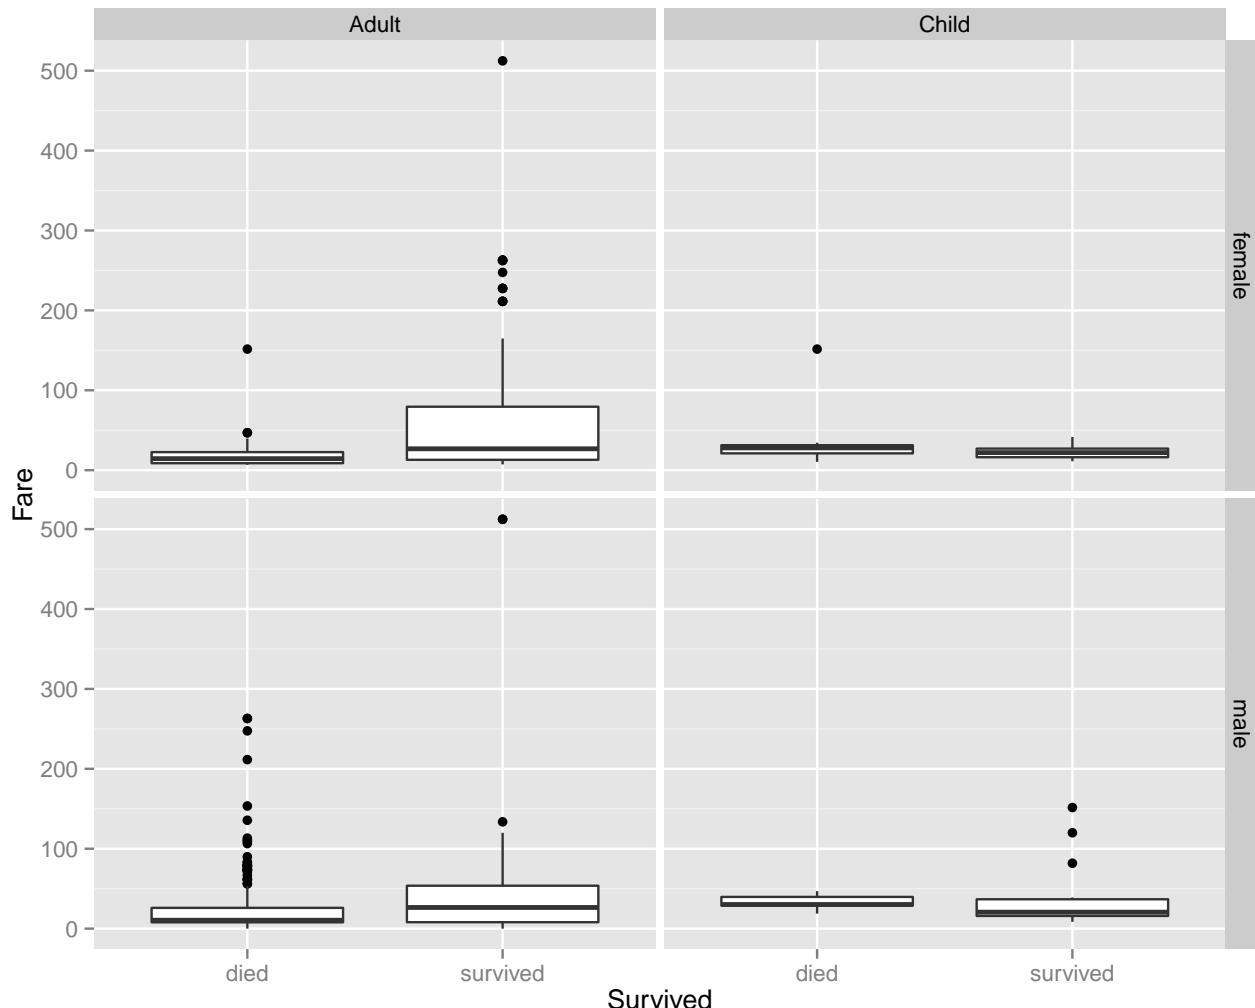
```

```

# Use labels instead of 0's and 1's
Child[Child==1] <- "Child"
Child[Child==2] <- "Adult"
# Appends the new column to the titanic dataset
titanic_with_child_column <- cbind(titanic, Child)
# Removes rows where age is NA
titanic_with_child_column <- titanic_with_child_column[!is.na(titanic_with_child_column$Child),]

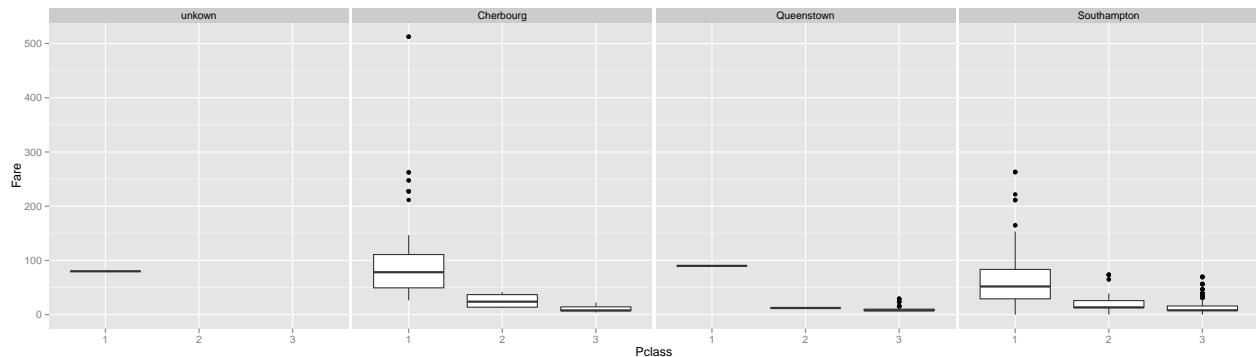
```

## Fare matters?



\*\*\*

## How about fare, ship class, port embarkation?



---

## Diamond data

---

### Overview of the diamond data

```
data(diamonds) # loading diamonds data set
head(diamonds, 16) # first few rows of diamond data set
```

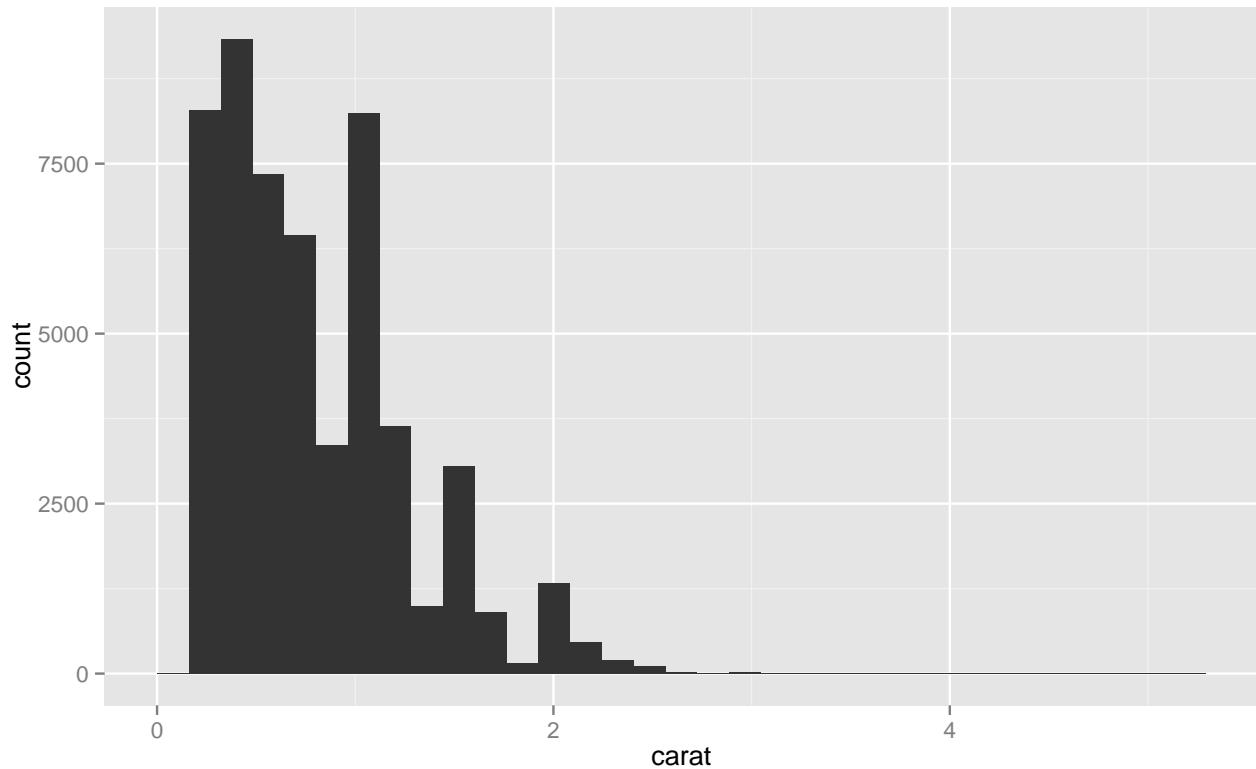
```
##   carat      cut color clarity depth table price     x     y     z
## 1  0.23    Ideal    E    SI2  61.5    55  326 3.95 3.98 2.43
## 2  0.21  Premium    E    SI1  59.8    61  326 3.89 3.84 2.31
## 3  0.23     Good    E    VS1  56.9    65  327 4.05 4.07 2.31
## 4  0.29  Premium    I    VS2  62.4    58  334 4.20 4.23 2.63
## 5  0.31     Good    J    SI2  63.3    58  335 4.34 4.35 2.75
## 6  0.24 Very Good    J   VVS2  62.8    57  336 3.94 3.96 2.48
## 7  0.24 Very Good    I   VVS1  62.3    57  336 3.95 3.98 2.47
## 8  0.26 Very Good    H    SI1  61.9    55  337 4.07 4.11 2.53
## 9  0.22      Fair    E    VS2  65.1    61  337 3.87 3.78 2.49
## 10 0.23 Very Good    H    VS1  59.4    61  338 4.00 4.05 2.39
## 11 0.30     Good    J    SI1  64.0    55  339 4.25 4.28 2.73
## 12 0.23    Ideal    J    VS1  62.8    56  340 3.93 3.90 2.46
## 13 0.22  Premium    F    SI1  60.4    61  342 3.88 3.84 2.33
## 14 0.31    Ideal    J    SI2  62.2    54  344 4.35 4.37 2.71
## 15 0.20  Premium    E    SI2  60.2    62  345 3.79 3.75 2.27
## 16 0.32  Premium    E     I1  60.9    58  345 4.38 4.42 2.68
```

---

### Histogram of carat

```
library(ggplot2)
ggplot(data=diamonds) + geom_histogram(aes(x=carat))

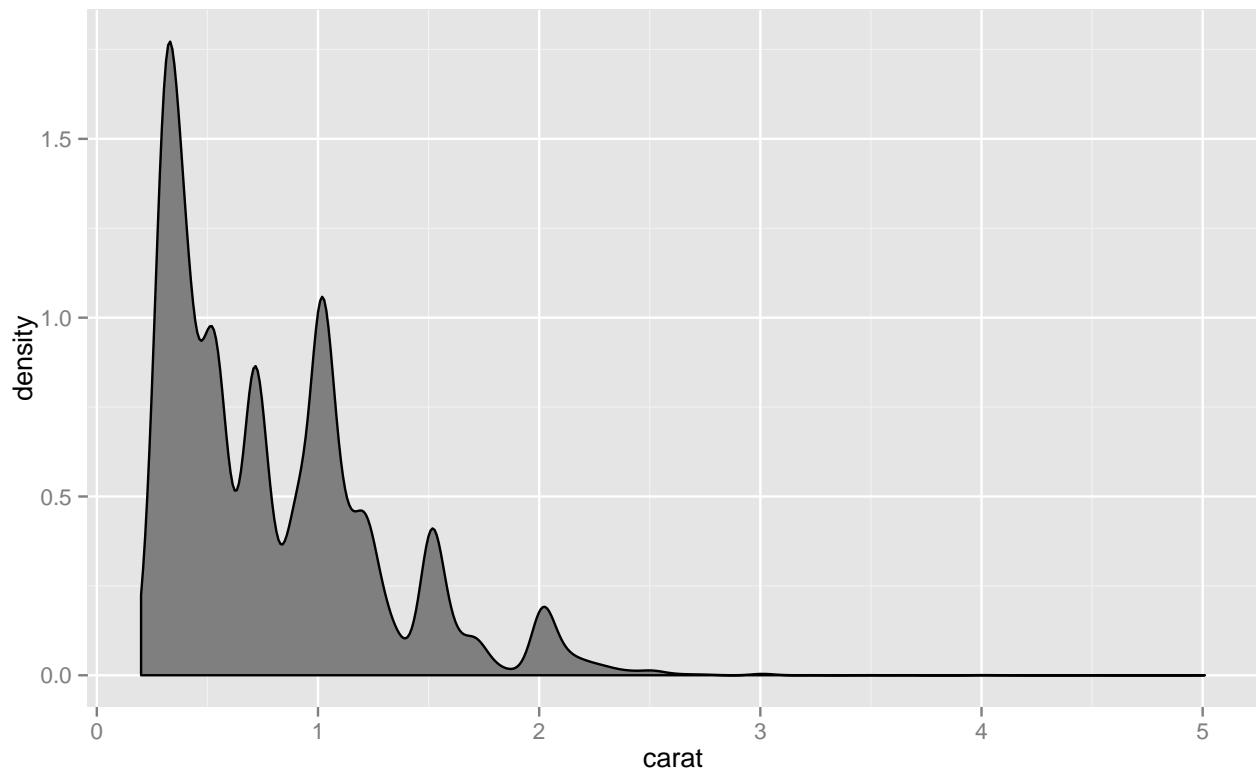
## stat_bin: binwidth defaulted to range/30. Use 'binwidth = x' to adjust this.
```



---

### Density plot of carat

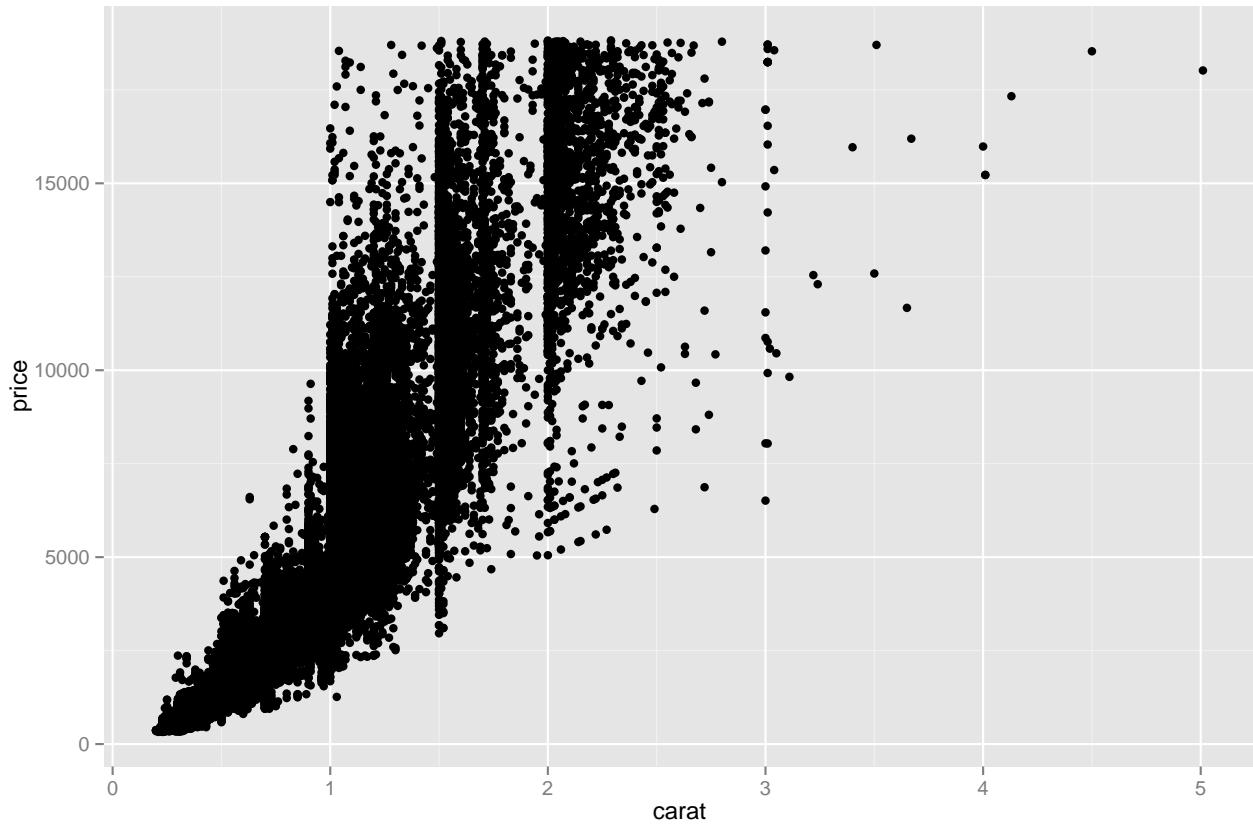
```
ggplot(data=diamonds) +
geom_density(aes(x=carat), fill="gray50")
```



---

### Scatter plots (carat vs. price)

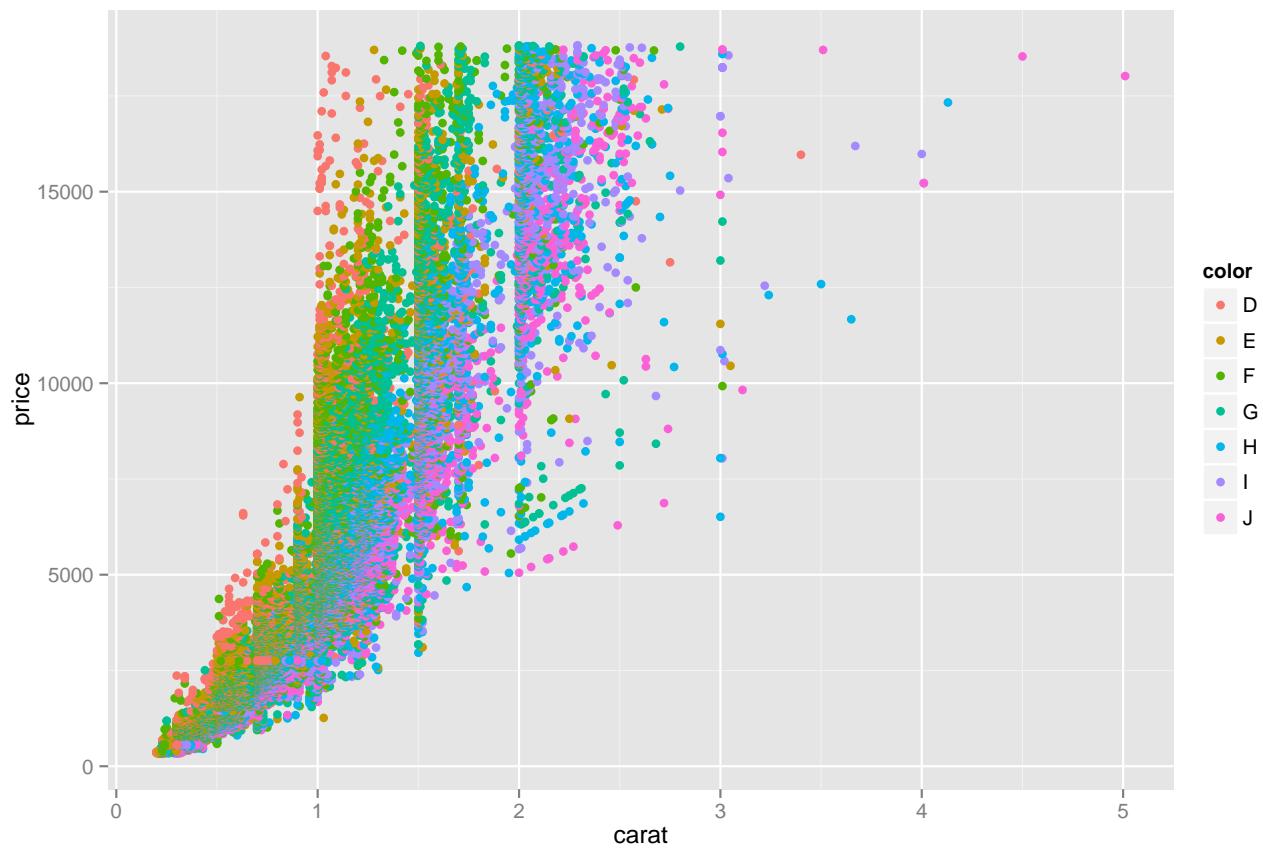
```
ggplot(diamonds, aes(x=carat,y=price)) + geom_point()
```



---

### Carat with colors

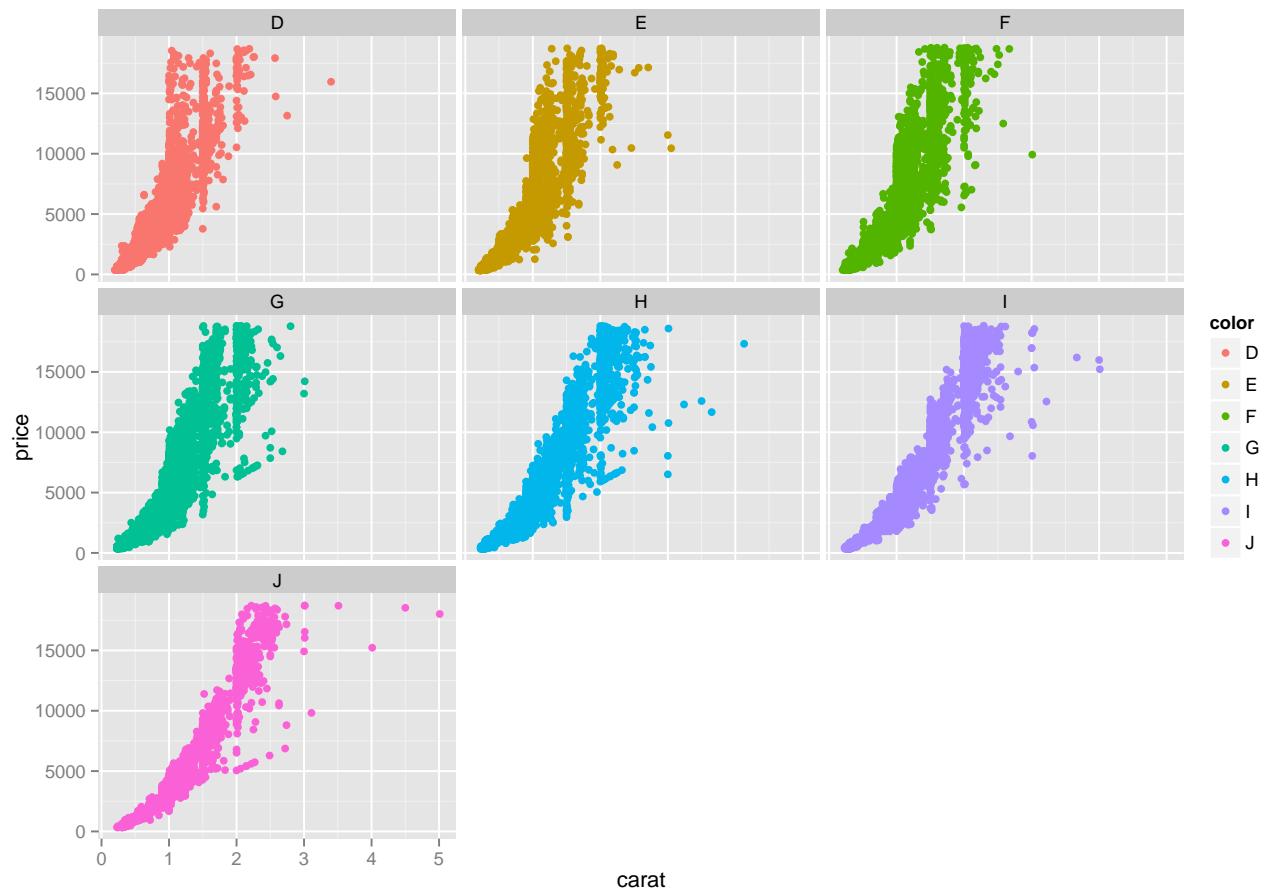
```
g = ggplot(diamonds, aes(x=carat, y=price)) # saving first layer as variable  
g + geom_point(aes(color=color)) # rendering first layer and adding another layer
```



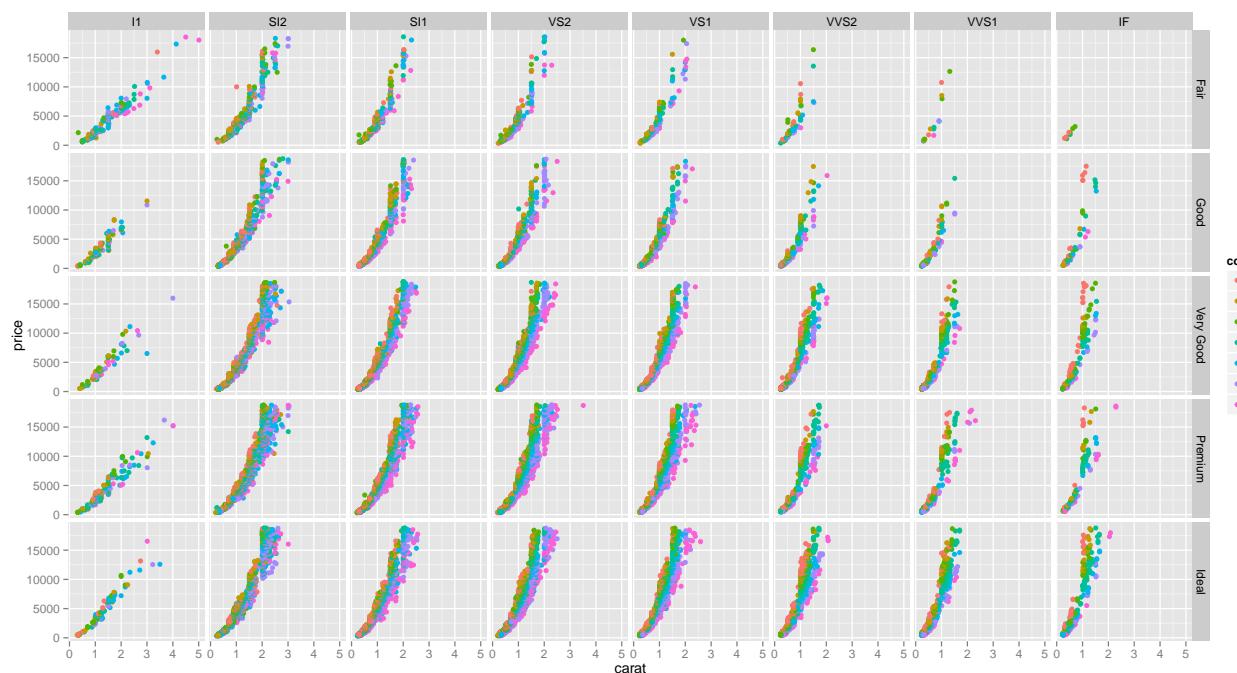
---

### Carat with colors (more details)

```
g + geom_point(aes(color=color)) + facet_wrap(~color)
```



Let's consider cut and clarity



---

## Your trun!

What is your knowledge of diamond's price after exploring this data?

---

## Interactive visualization in R - rCharts

- What is rCharts?  
Is an R package to create, customize and publish interactive javascript visualizations from R using a familiar lattice style plotting interface.
  - What rCharts can make and how?  
Quick start at: <http://ramnathv.github.io/rCharts/>
  - A list of interactive visualization in R can be found at:  
<http://ouzor.github.io/blog/2014/11/21/interactive-visualizations.html>
- 

## Tell your story - R Markdown

- R Markdown is an authoring format that enables easy creation of dynamic documents, presentations, and reports from R.
  - It combines the core syntax of markdown (an easy-to-write plain text format) with embedded R code chunks that are run so their output can be included in the final document.
  - Many available output formats including HTML, PDF, and MS Word
  - **Installation**  
Use RStudio: already installed  
Outside of RStudio: `install.packages("rmarkdown")`. A recent version of pandoc ( $\geq 1.12.3$ ) is also required. See <https://github.com/rstudio/rmarkdown/blob/master/PANDOC.MD> to install it.
- 

## Check out Markdown first

Markdown is a markup language with plain text formatting syntax designed so that it can be converted to HTML and many other formats using a tool by the same name.

One minute you get the point, and always check the cheat sheets  
<https://github.com/adam-p/markdown-here/wiki/Markdown-Cheatsheet#lists>

---

## Then, R Markdown sample code

Download the template:

[https://github.com/datasets/blob/master/rmarkdown\\_template.Rmd](https://github.com/datasets/blob/master/rmarkdown_template.Rmd)

## R Markdown

- YAML header
- Edit Markdown, and R chunks
- Run!
  - RStudio: knitr button
  - Command line: render("file.Rmd")

Cheat sheet of rmarkdown:

<http://www.rstudio.com/wp-content/uploads/2015/02/rmarkdown-cheatsheet.pdf>

---

## Present your story of Titanic!

Use \* Titanic data \* Plotting functions in R \* R Markdown template \* **The heart of data explorer** to write your story of Titanic...

---

Hope this is inspiring :)

[Titanic](#)